

CAFA-IVR: Counterfactual ASR Failure Attribution for Conversational IVR Testing

Vijay Kumar Sridharan
IEEE Senior Member
Frisco, TX, USA
sridharanvijaykumar@gmail.com

Abstract—Automatic speech recognition (ASR) quality in conversational interactive voice response (IVR) systems is commonly summarized by word error rate (WER), although a production voice journey ultimately succeeds or fails at a downstream semantic task. We introduce CAFA-IVR, a vendor-neutral counterfactual testing framework that executes each semantic case through two paths: reference text $\rightarrow$ downstream system and audio $\rightarrow$ ASR $\rightarrow$ the same downstream system. A failure that appears only after the speech path is introduced is quantified by ASR-Attributable Intent Failure Rate (ASR-IFR). CAFA-IVR also defines a portable attribution matrix, critical-entity scoring, and a CI/CD integration contract. We validate the method with 540 locally executed ASR trials: 180 PocketSphinx decodes and 360 Conformer-CTC neural decodes across clean, telephony, noise, and speaking-rate conditions. In the neural study, clean WER was 11.5%, intent accuracy 70.0%, and ASR-IFR 10.0%; under 5-dB noise these became 122.9%, 10.0%, and 71.7%. We additionally map the framework to a public human banking-speech corpus to establish external domain/channel applicability while keeping corpus-author results separate from our measurements. CAFA-IVR is released as a reference specification and scorer so organizations can reproduce attribution, regression gates, and evidence retention without adopting a particular speech vendor.

Index Terms—automatic speech recognition, conversational AI, IVR, testing, failure attribution, word error rate, spoken language understanding

I. INTRODUCTION

A conversational IVR is a layered system: audio capture and channel processing feed ASR, whose transcript is interpreted by natural-language understanding (NLU), an agent, or dialogue logic. When the observed intent or task is wrong, the final failure alone does not identify whether the speech layer or the downstream system caused it.

WER remains the standard lexical diagnostic,

$$\text{WER} = \frac{S + D + I}{N}, \quad (1)$$

where S , D , and I are substitutions, deletions, and insertions relative to N reference words. However, lexical distance is not an attribution mechanism. Semantic-ASR research has similarly motivated alternatives or complements to pure lexical scoring [1], [2], while spoken-language-understanding studies document downstream degradation under ASR noise [3]. Audio/text bot workbenches demonstrate the practicality of multimodal regression testing [4]. The remaining operational question is: *did the addition of the speech path cause an otherwise valid semantic test to fail?*

Reference text $r_i \rightarrow$ Same downstream system $\rightarrow T_i$
Audio $a_i \rightarrow$ ASR $\rightarrow$ Same downstream system $\rightarrow A_i$
(T_i, A_i) $\rightarrow$ Attribution state $\rightarrow$ WER / CEER / ASR-IFR

Fig. 1. CAFA-IVR paired counterfactual execution. Only the audio path introduces ASR; downstream components and the task oracle remain frozen.

TABLE I
CAFA-IVR ATTRIBUTION MATRIX.

Text	Audio	Attribution
Pass	Pass	HEALTHY
Pass	Fail	SPEECH_ATTRIBUTABLE
Fail	Fail	DOWNSTREAM_OR_TEST
Fail	Pass	CONTEXT_OR_ORACLE

CAFA-IVR makes four contributions: (1) a paired text/audio counterfactual; (2) ASR-IFR; (3) a portable attribution and critical-entity contract suitable for CI/CD; and (4) locally executed real-decoding pilots plus an external human-banking-corpus mapping. The reference scorer, specification, examples, and empirical artifacts are openly available at <https://github.com/sridharanvijaykumar/cafa-ivr>. The released artifact (v1.0.x, Apache-2.0) reproduces metric computation from the supplied decoded outputs. Audio generation, ASR decoding, perturbation, and bootstrap resampling are not included in this release; reported decode-level and bootstrap results are therefore not reproducible end to end from the artifact alone.

II. CAFA-IVR FRAMEWORK

For semantic case i , let r_i be the reference text and a_i its audio realization. CAFA-IVR executes

$$T_i : r_i \rightarrow \text{Downstream}, \quad (2)$$

$$A_i : a_i \rightarrow \text{ASR} \rightarrow \text{Downstream}. \quad (3)$$

The downstream model, prompt, tools, dialogue state, and test oracle should be frozen between the two paths. Let $T_i = 1$ and $A_i = 1$ denote successful semantic outcomes. Table I defines the attribution states.

Define $F_i = \neg(T_i = 1 \wedge A_i = 0)$. The primary metric is

$$\text{ASR-IFR} = \frac{1}{N} \sum_{i=1}^N F_i. \quad (4)$$

Thus ASR-IFR measures the fraction of the complete paired suite converted from a downstream-capable reference-text

success to an audio-path failure. For teams that want risk conditioned only on eligible controls, CAFA-IVR additionally reports

$$\text{cASR-IFR} = \frac{\sum_i F_i}{\sum_i T_i}, \quad (5)$$

while keeping the all-trial ASR-IFR as the primary release metric.

CAFA-IVR also retains WER and defines Critical-Entity Error Rate (CEER) over governed values such as amounts, dates, negation, authentication values, account types, and high-risk action verbs. If C_i and $\hat{C}_i$ are normalized expected and observed critical-entity multisets,

$$\text{CEER} = \frac{\sum_i (|C_i \setminus \hat{C}_i| + |\hat{C}_i \setminus C_i|)}{\sum_i |C_i|}. \quad (6)$$

An optional realized-consequence score (CIER) may be added. Let $\mathcal{I}$ be the set of trials carrying a recognized impact level, and let $w(L_i)$ be the deployment-selected weight for level L_i . Then

$$\text{CIER} = \frac{1}{|\mathcal{I}|} \sum_{i \in \mathcal{I}} (1 - A_i) w(L_i). \quad (7)$$

Successful audio-path tasks contribute zero rather than being excluded, while failed tasks contribute their impact weight; CIER is therefore a severity-weighted failure rate over the impact-labelled population rather than an average severity of failures. It counts audio-path failures regardless of attribution. Speech-attributable consequence is obtained by restricting the population to the `SPEECH_ATTRIBUTABLE` state. The severity weights remain deployment policy rather than universal constants.

III. PORTABLE TEST AND RELEASE CONTRACT

The minimum trial record contains a stable test ID, acoustic condition, reference text, ASR transcript, expected/observed semantic outcome, text-control result, audio-path task result, and engine/version metadata. Critical entities and risk tier are recommended for high-consequence journeys.

The framework is deliberately vendor-neutral. The downstream component can be a classical intent model, deterministic workflow, LLM agent, or tool-calling system, provided a reproducible task oracle exists. The operational algorithm is:

- 1) freeze candidate ASR and downstream versions;
- 2) run all reference-text controls;
- 3) run the identical cases through audio and ASR;
- 4) compute WER, ASR-IFR, and CEER where applicable;
- 5) compare with the last approved baseline;
- 6) route failures using Table I;
- 7) retain failed audio, transcript, versions, and outcome for replay.

CAFA-IVR does not prescribe a universal numeric release threshold. Instead, teams can gate on regression relative to an approved baseline, with stricter policy for high-risk intents. For example, an organization may require no unexplained ASR-IFR increase and no new critical-entity errors on governed payment, fraud, or card-control journeys. WER remains a diagnostic but is not the sole release decision.

TABLE II
LOCALLY EXECUTED CAFA-IVR STUDIES.

Study	Trials	Text ctrl. Acc. (%)	WER (%)	ASR-IFR (%)
PocketSphinx	180	96.7	88.9	75.0
Conformer-CTC	360	80.0	64.8	39.4

IV. EXPERIMENTAL DESIGN

A. Legacy Pilot

A 600-utterance benchmark spans 30 IVR intents. One utterance per intent was held out and the other 570 trained a TF-IDF/logistic-regression downstream classifier. Held-out text-control accuracy was 96.7%. eSpeak generated audio under six deterministic conditions: clean, 8-kHz telephony, 15-dB noise, 5-dB noise, 1.15× rate, and 8-kHz plus 15-dB noise. PocketSphinx produced 180 actual ASR decodes. Overall WER was 88.9%, intent accuracy 22.2%, and ASR-IFR 75.0%. This weak recognizer/synthetic-speech combination is retained as pipeline validation rather than a modern-ASR performance claim.

B. Neural Pilot

A second experiment used a 60-phrase closed IVR grammar (two phrases per intent) synthesized by ten eSpeak English voices. V01–V08 (480 clips) trained the acoustic model, V09 (60) was validation, and unseen V10 (60) was the test speaker. The acoustic model used 64-bin log-Mel features, a 96-dimensional representation, three Conformer blocks, four attention heads, a 256-dimensional feed-forward layer, kernel size 15, and character CTC output (516,029 parameters). Conformer combines convolutional local modeling with Transformer-style global modeling [5]. The frozen epoch-20 model used a constrained 60-phrase CTC decoder. The downstream NLU was separately trained on the other 540 benchmark texts and obtained 80.0% text-control accuracy on these 60 cases.

The V10 test audio was evaluated under the same six conditions, producing 360 actual neural decodes with no condition-specific retraining.

V. RESULTS

Table II summarizes the two locally executed studies. They validate the same attribution pipeline under different recognizer assumptions and are not presented as an apples-to-apples ASR leaderboard.

Table III reports measured neural results. Clean WER was 11.5%, exact phrase recognition 88.3%, downstream intent accuracy 70.0%, and ASR-IFR 10.0%. Telephony increased ASR-IFR to 20.0%, and severe 5-dB noise increased it to 71.7%. WER can exceed 100% when total word edits exceed the reference length.

Nonparametric trial-level bootstrap intervals used 5,000 resamples with a fixed seed. Bootstrap 95% intervals for clean speech were 4.2–20.1% WER, 58.3–81.7% intent accuracy,

TABLE III
MEASURED CONFORMER-CTC RESULTS, 60 TRIALS/CONDITION.

Condition	WER (%)	Phrase	Intent	IFR (%)
		Exact (%)	Acc. (%)	
Clean	11.5	88.3	70.0	10.0
8 kHz	27.2	68.3	60.0	20.0
15 dB	74.1	41.7	45.0	36.7
5 dB	122.9	6.7	10.0	71.7
Rate 1.15×	41.2	55.0	48.3	31.7
8 kHz + 15 dB	112.1	18.3	13.3	66.7
Overall	64.8	46.4	41.1	39.4

and 3.3–18.3% ASR-IFR. Across all 360 trials, the corresponding intervals were 57.7–72.1%, 35.8–46.1%, and 34.4–44.4%. The neural text control passed in 288 of 360 paired trials; 142 were speech-attributable failures, yielding 39.4% all-trial ASR-IFR and 49.3% conditional cASR-IFR.

The distinction from WER is visible at trial level. “my physical card is missing” became “the card is missing” (WER 0.40) while preserving `lost_card`. By contrast, “what are my latest payments” became “what are my last transactions” (WER 0.40), changing `payment_history` to `recent_transactions`. Equal WER therefore need not imply equal semantic consequence.

The text counterfactual additionally prevented false ASR attribution. Twelve of the 60 neural reference texts failed the downstream NLU before speech was introduced. Across six conditions, 34 trials had an exact ASR transcript but belonged to those text-control failures; a voice-only dashboard could otherwise direct debugging toward ASR.

The pilot reported here exercises ASR-IFR and WER across six acoustic conditions: clean, telephony at 8 kHz, additive noise at 15 dB and 5 dB, rate 1.15×, and telephony with noise. It does not exercise CEER or CIER. The 360-trial set contains no critical-entity annotations and no impact labels, so both metrics are specified and implemented but not empirically demonstrated in this release. Their evaluation is left to an entity- and impact-annotated subset in future work.

VI. EXTERNAL HUMAN-SPEECH MAPPING

The local studies use synthetic speech, so we map the protocol to HarperValleyBank (HVB), a public domain-specific consumer-banking dialogue corpus containing 23.7 hours, 1,446 conversations, 25,730 utterances, 59 speakers, eight task classes, and separate 8-kHz call sides [6]. An online SLUE-HVB sample was frozen across six banking intents with speaker IDs, segment IDs, timestamps, human transcripts, and license/source metadata [7]. This is an external source-validation layer, not a third locally executed ASR experiment: the online audio binaries could not be materialized in the execution environment, so no sample-wide WER or ASR-IFR is claimed. HVB nevertheless supplies the human telephony channel, intent labels, and corrected transcript structure needed by CAFA-IVR without changing the test contract.

TABLE IV
MINIMUM ADOPTER INTERFACE.

Artifact	Purpose
Reference text	Counterfactual control input
Audio + transcript	Speech-path input and lexical diagnostic
Task oracle	Frozen success/failure decision for both paths
Engine metadata	Reproducible ASR/downstream versioning
Baseline summary	Regression comparison in CI/CD

VII. INDUSTRIAL DEPLOYMENT

The research contribution is packaged as a versioned reference specification and scorer rather than only a paper. An adopter supplies trial-level outputs from its existing ASR/NLU stack; the scorer produces attributed trials, condition summaries, and baseline/candidate comparisons. A CI job can fail when organization-defined ASR-IFR or CEER regression limits are crossed. Because the same test schema applies to conventional NLU and generative agents, the framework can be inserted around an existing voice stack without replacing its inference components.

Independent implementation should record the CAFA-IVR version, ASR/downstream versions, test corpus, thresholds, resulting engineering decision, and retained evidence. This makes both successful adoption and negative evaluations independently auditable rather than dependent on the original author. The open implementation exposes `score` and `compare` commands and adapter examples so a team can adopt the attribution contract without changing its ASR vendor or dialogue platform.

VIII. LIMITATIONS

Both locally executed studies use synthetic speech. The neural decoder is closed-set and phrase-overlapping across training/test speakers, so it is an in-domain constrained recognizer rather than a pretrained open-vocabulary benchmark. The 80% neural text-control accuracy also limits eligible attribution cases. Human-speech multi-engine inference, stronger frozen downstream controls, validated entity normalization, and speaker-level paired statistics are needed before making broader robustness claims.

IX. CONCLUSION

CAFA-IVR converts conversational voice regression testing from a single transcript-quality number into a paired attribution procedure. Across 540 actual local ASR decodes, it separates lexical mismatch from failures that appear only after speech is introduced and avoids blaming ASR for downstream cases that already fail on reference text. The framework is deliberately implementable as a lightweight layer around existing voice systems, with WER, ASR-IFR, CEER, attribution states, reproducibility metadata, and baseline release gates exposed through a portable specification and reference scorer.

REFERENCES

- [1] S. Kim et al., “Semantic Distance: A New Metric for ASR Performance Analysis Towards Spoken Language Understanding,” arXiv:2104.02138, 2021.

- [2] S. Roy, “Semantic-WER: A Unified Metric for the Evaluation of ASR Transcript for End Usability,” arXiv:2106.02016, 2021.
- [3] S. Sengupta, J. Krone, and S. Mansour, “On the Robustness of Intent Classification and Slot Labeling in Goal-oriented Dialog Systems to Real-world Noise,” in *Proc. NLP4ConvAI*, 2021, pp. 68–79.
- [4] Amazon Web Services, “Evaluating Amazon Lex V2 bot performance with Test Workbench,” *Amazon Lex V2 Developer Guide*, accessed Aug. 2026.
- [5] A. Gulati et al., “Conformer: Convolution-augmented Transformer for Speech Recognition,” in *Proc. Interspeech*, 2020, pp. 5036–5040.
- [6] M. Wu, J. Nafziger, A. Scodary, and A. Maas, “HarperValleyBank: A Domain-Specific Spoken Dialog Corpus,” arXiv:2010.13929, 2020.
- [7] S. Shon et al., “SLUE Phase-2: A Benchmark Suite of Diverse Spoken Language Understanding Tasks,” arXiv:2212.10525, 2022.